

## **Lecture 2**

# **Relationships between Categorical Variables**

- ① Review
- ② Two (or More) Categorical Variables
- ③ Proportions and Probabilities
- ④ Joint and Conditional Distributions

- 1 Review
- 2 Two (or More) Categorical Variables
- 3 Proportions and Probabilities
- 4 Joint and Conditional Distributions

# Types of Data

Type	On Disk	In Python
tabular		

---

# Types of Data

Type	On Disk	In Python
tabular	CSV	

# Types of Data

Type	On Disk	In Python
tabular	CSV	DataFrame

---

# Types of Data

Type	On Disk	In Python
tabular	CSV	DataFrame
hierarchical		

---

# Types of Data

Type	On Disk	In Python
tabular	CSV	DataFrame
hierarchical	JSON	

---

# Types of Data

Type	On Disk	In Python
tabular	CSV	DataFrame
hierarchical	JSON	dict

```
{  
  "name": "Alice",  
  "age": 28,  
  "married": false,  
  "children": null,  
  "scores": [90, 85, 88],  
  "address": {  
    "city": "London",  
    "zip": 12345  
  }  
}
```

# Types of Data

Type	On Disk	In Python
tabular	CSV	DataFrame
hierarchical	JSON	dict
textual		

---

# Types of Data

Type	On Disk	In Python
tabular	CSV	DataFrame
hierarchical	JSON	dict
textual	plaintext	

---

# Types of Data

Type	On Disk	In Python
tabular	CSV	DataFrame
hierarchical	JSON	dict
textual	plaintext	string

# Types of Data

Type	On Disk	In Python
tabular	CSV	DataFrame
hierarchical	JSON	dict
textual	plaintext	string
geospatial	???	???

```
{  
  "type": "Feature",  
  "geometry": {  
    "type": "Point",  
    "coordinates": [-0.1276, 51.5072]  
  },  
  "properties": {  
    "name": "London",  
    "country": "United Kingdom",  
    "population": "8,982,000"  
  }  
}
```

- Vector data: .shp, .geojson, .kml, .gpx, .gpkg
- Raster data: .tif, .img, .asc
- Databases / Mixed: .gdb, .gpkg

Example for geojson file.

# Types of Variables

```
import pandas as pd  
df = pd.read_csv("https://datasci112.stanford.edu/data/titanic.csv")  
df
```

```
pd.read_json("file.json") or json.load()
```

pandas / json

# Types of Variables

```
import pandas as pd  
df = pd.read_csv("https://datasci112.stanford.edu/data/titanic.csv")  
df
```

		name	pclass	survived	sex	age	sibsp	parch	ticket	fare	cabin	embarked	boat	body	home.dest
0		Allen, Miss. Elisabeth Walton	1	1	female	29.0000	0	0	24160	211.3375	B5	S	2	NaN	St Louis, MO
1		Allison, Master. Hudson Trevor	1	1	male	0.9167	1	2	113781	151.5500	C22 C26	S	11	NaN	Montreal, PQ / Chesterville, ON
2		Allison, Miss. Helen Loraine	1	0	female	2.0000	1	2	113781	151.5500	C22 C26	S	NaN	NaN	Montreal, PQ / Chesterville, ON
3		Allison, Mr. Hudson Joshua Creighton	1	0	male	30.0000	1	2	113781	151.5500	C22 C26	S	NaN	135.0	Montreal, PQ / Chesterville, ON
4		Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	1	0	female	25.0000	1	2	113781	151.5500	C22 C26	S	NaN	NaN	Montreal, PQ / Chesterville, ON
...		...	...	...	...	...	...	...	...	...	...	...	...	...	...
1304		Zabour, Miss. Hileni	3	0	female	14.5000	1	0	2665	14.4542	NaN	C	NaN	328.0	NaN
1305		Zabour, Miss. Thamrine	3	0	female	NaN	1	0	2665	14.4542	NaN	C	NaN	NaN	NaN
1306		Zakarian, Mr. Mapriededer	3	0	male	26.5000	0	0	2656	7.2250	NaN	C	NaN	304.0	NaN
1307		Zakarian, Mr. Ortin	3	0	male	27.0000	0	0	2670	7.2250	NaN	C	NaN	NaN	NaN
1308		Zimmerman, Mr. Leo	3	0	male	29.0000	0	0	315082	7.8750	NaN	S	NaN	NaN	NaN

1309 rows x 14 columns

# Types of Variables

```
import pandas as pd  
df = pd.read_csv("https://datasci112.stanford.edu/data/titanic.csv")  
df
```

**variables**

The diagram shows a DataFrame with 1309 rows and 14 columns. The columns are labeled: name, pclass, survived, sex, age, sibsp, parch, ticket, fare, cabin, embarked, boat, body, and home.dest. A vertical stack of arrows on the left points from the row index 0 up to 1308, with the label 'observational units' written vertically next to it. Another set of arrows points from the column headers to specific rows: one arrow from 'name' to row 0, another from 'age' to row 0, and a third from 'cabin' to row 0. Subsequent arrows point from each of these three columns to rows 1 through 4.

		name	pclass	survived	sex	age	sibsp	parch	ticket	fare	cabin	embarked	boat	body	home.dest
0		Allen, Miss. Elisabeth Walton	1	1	female	29.0000	0	0	24160	211.3375	B5	S	2	NaN	St Louis, MO
1		Allison, Master. Hudson Trevor	1	1	male	0.9167	1	2	113781	151.5500	C22 C26	S	11	NaN	Montreal, PQ / Chesterville, ON
2		Allison, Miss. Helen Loraine	1	0	female	2.0000	1	2	113781	151.5500	C22 C26	S	NaN	NaN	Montreal, PQ / Chesterville, ON
3		Allison, Mr. Hudson Joshua Creighton	1	0	male	30.0000	1	2	113781	151.5500	C22 C26	S	NaN	135.0	Montreal, PQ / Chesterville, ON
4		Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	1	0	female	25.0000	1	2	113781	151.5500	C22 C26	S	NaN	NaN	Montreal, PQ / Chesterville, ON
...		...	...	...	...	...	...	...	...	...	...	...	...	...	
1304		Zabour, Miss. Hileni	3	0	female	14.5000	1	0	2665	14.4542	NaN	C	NaN	328.0	NaN
1305		Zabour, Miss. Thamrine	3	0	female	NaN	1	0	2665	14.4542	NaN	C	NaN	NaN	NaN
1306		Zakarian, Mr. Mapriededer	3	0	male	26.5000	0	0	2656	7.2250	NaN	C	NaN	304.0	NaN
1307		Zakarian, Mr. Ortin	3	0	male	27.0000	0	0	2670	7.2250	NaN	C	NaN	NaN	NaN
1308		Zimmerman, Mr. Leo	3	0	male	29.0000	0	0	315082	7.8750	NaN	S	NaN	NaN	NaN

1309 rows x 14 columns

# Types of Variables

```
import pandas as pd  
df = pd.read_csv("https://datasci112.stanford.edu/data/titanic.csv")  
df
```

**variables**

		name	pclass	survived	sex	age	sibsp	parch	ticket	fare	cabin	embarked	boat	body	home.dest
0		Allen, Miss. Elisabeth Walton	1	1	female	29.0000	0	0	24160	211.3375	B5	S	2	NaN	St Louis, MO
1		Allison, Master. Hudson Trevor	1	1	male	0.9167	1	2	113781	151.5500	C22 C26	S	11	NaN	Montreal, PQ / Chesterville, ON
2		Allison, Miss. Helen Loraine	1	0	female	2.0000	1	2	113781	151.5500	C22 C26	S	NaN	NaN	Montreal, PQ / Chesterville, ON
3		Allison, Mr. Hudson Joshua Creighton	1	0	male	30.0000	1	2	113781	151.5500	C22 C26	S	NaN	135.0	Montreal, PQ / Chesterville, ON
4		Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	1	0	female	25.0000	1	2	113781	151.5500	C22 C26	S	NaN	NaN	Montreal, PQ / Chesterville, ON
...		...	...	...	...	...	...	...	...	...	...	...	...	...	...
1304		Zabour, Miss. Hileni	3	0	female	14.5000	1	0	2665	14.4542	NaN	C	NaN	328.0	NaN
1305		Zabour, Miss. Thamrine	3	0	female	NaN	1	0	2665	14.4542	NaN	C	NaN	NaN	NaN
1306		Zakarian, Mr. Mapriededer	3	0	male	26.5000	0	0	2656	7.2250	NaN	C	NaN	304.0	NaN
1307		Zakarian, Mr. Ortin	3	0	male	27.0000	0	0	2670	7.2250	NaN	C	NaN	NaN	NaN
1308		Zimmerman, Mr. Leo	3	0	male	29.0000	0	0	315082	7.8750	NaN	S	NaN	NaN	NaN

**observational units**

**quantitative variables**

**categorical variables**

# One Categorical Variable

To *summarize* a categorical variable, we report the **counts** of each possible category.

```
df["pclass"].value_counts().sort_index()
```

```
1    323
2    277
3    709
Name: pclass, dtype: int64
```

# One Categorical Variable

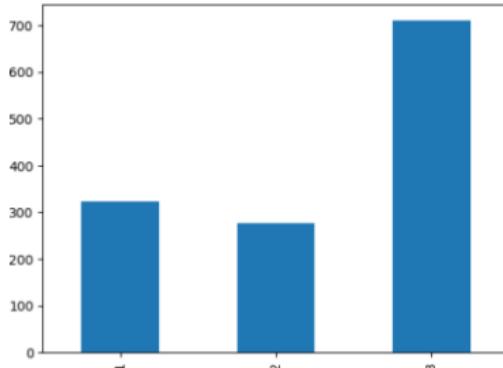
To *summarize* a categorical variable, we report the **counts** of each possible category.

```
df["pclass"].value_counts().sort_index()
```

```
1    323
2    277
3    709
Name: pclass, dtype: int64
```

To *visualize* a categorical variable, we make a **bar plot**.

```
df["pclass"].value_counts().sort_index().plot.bar()
```



- 1 Review
- 2 Two (or More) Categorical Variables
- 3 Proportions and Probabilities
- 4 Joint and Conditional Distributions

# Selecting Columns

Note that we selected a single column by passing the column name as a key to the DataFrame.

```
df["pclass"]
```

```
0      1  
1      1  
2      1  
..  
1306    3  
1307    3  
1308    3  
Name: pclass, Length: 1309, dtype: int64
```

# Selecting Columns

Note that we selected a single column by passing the column name as a key to the DataFrame.

```
df["pclass"]
```

```
0      1  
1      1  
2      1  
..  
1306    3  
1307    3  
1308    3  
  
Name: pclass, Length: 1309, dtype: int64
```

The result is a one-dimensional pandas object called a Series.

# Selecting Columns

Note that we selected a single column by passing the column name as a key to the DataFrame.

```
df["pclass"]
```

```
0      1  
1      1  
2      1  
..  
1306    3  
1307    3  
1308    3  
  
Name: pclass, Length: 1309, dtype: int64
```

The result is a one-dimensional pandas object called a Series.

We can select multiple columns by passing a [list](#) of column names.



# Selecting Columns

Note that we selected a single column by passing the column name as a key to the DataFrame.

```
df["pclass"]
```

```
0      1  
1      1  
2      1  
..  
1306    3  
1307    3  
1308    3  
  
Name: pclass, Length: 1309, dtype: int64
```

The result is a one-dimensional pandas object called a Series.

We can select multiple columns by passing a [list](#) of column names.

```
df[["pclass", "survived"]]
```

# Selecting Columns

Note that we selected a single column by passing the column name as a key to the DataFrame.

```
df["pclass"]
```

```
0      1  
1      1  
2      1  
..  
1306    3  
1307    3  
1308    3  
  
Name: pclass, Length: 1309, dtype: int64
```

The result is a one-dimensional pandas object called a Series.

We can select multiple columns by passing a [list](#) of column names.

```
df[["pclass", "survived"]]
```

```
   pclass  survived  
0      1         1  
1      1         1  
2      1         0  
3      1         0  
4      1         0  
..    ...       ...  
1304    3         0  
1305    3         0  
1306    3         0  
1307    3         0  
1308    3         0  
  
1309 rows x 2 columns
```

The result is two-dimensional, another smaller DataFrame.

# Selecting Columns

Note that we selected a single column by passing the column name as a key to the DataFrame.

```
df["pclass"]
```

```
0      1  
1      1  
2      1  
..  
1306    3  
1307    3  
1308    3  
  
Name: pclass, Length: 1309, dtype: int64
```

The result is a one-dimensional pandas object called a Series.

We can select multiple columns by passing a [list](#) of column names.

```
df[["pclass", "survived"]]
```

```
   pclass  survived  
0      1          1  
1      1          1  
2      1          0  
3      1          0  
4      1          0  
..    ...        ...  
1304    3          0  
1305    3          0  
1306    3          0  
1307    3          0  
1308    3          0  
  
1309 rows x 2 columns
```

The result is two-dimensional, another smaller DataFrame.

How do we make sense of multiple variables at once?

## Summarizing Multiple Categorical Variables

To summarize multiple categorical variables, we report the **counts** of every possible combination of categories.

## Summarizing Multiple Categorical Variables

To summarize multiple categorical variables, we report the **counts** of every possible combination of categories.

We can use the `.value_counts()` method of DataFrame.

# Summarizing Multiple Categorical Variables

To summarize multiple categorical variables, we report the **counts** of every possible combination of categories.

We can use the `.value_counts()` method of DataFrame.

```
df[["pclass", "survived"]].value_counts()
```

# Summarizing Multiple Categorical Variables

To summarize multiple categorical variables, we report the **counts** of every possible combination of categories.

We can use the `.value_counts()` method of DataFrame.

```
df[["pclass", "survived"]].value_counts()
```

```
pclass    survived
3          0            528
1          1            200
3          1            181
2          0            158
1          0            123
2          1            119
dtype: int64
```

# Summarizing Multiple Categorical Variables

To summarize multiple categorical variables, we report the **counts** of every possible combination of categories.

We can use the `.value_counts()` method of DataFrame.

```
df[["pclass", "survived"]].value_counts()
```

```
pclass    survived
3          0            528
1          1            200
3          1            181
2          0            158
1          0            123
2          1            119
dtype: int64
```

Note that the result is a Series, with a multi-level index, one for each variable!

# Summarizing Multiple Categorical Variables

```
pclass    survived  
3          0            528  
1          1            200  
3          1            181  
2          0            158  
1          0            123  
2          1            119  
dtype: int64
```

Let's make this information easier to read by arranging one variable along the rows and the other along the columns.

# Summarizing Multiple Categorical Variables

```
pclass    survived  
3          0            528  
1          1            200  
3          1            181  
2          0            158  
1          0            123  
2          1            119  
dtype: int64
```

Let's make this information easier to read by arranging one variable along the rows and the other along the columns.

```
(df[["pclass", "survived"]].value_counts().  
unstack())
```

# Summarizing Multiple Categorical Variables

```
pclass    survived
3          0            528
1          1            200
3          1            181
2          0            158
1          0            123
2          1            119
dtype: int64
```

Let's make this information easier to read by arranging one variable along the rows and the other along the columns.

```
(df[["pclass", "survived"]].value_counts().  
unstack())
```

```
survived    0    1
pclass
1          123  200
2          158  119
3          528  181
```

# Summarizing Multiple Categorical Variables

```
pclass    survived
3          0            528
1          1            200
3          1            181
2          0            158
1          0            123
2          1            119
dtype: int64
```

Let's make this information easier to read by arranging one variable along the rows and the other along the columns.

```
(df[["pclass", "survived"]].value_counts().  
unstack())
```

survived	0	1
pclass		
1	123	200
2	158	119
3	528	181

This representation is called a **two-way table** or a **crosstab** (short for “cross-tabulation”).

# Visualizing Multiple Categorical Variables

# Visualizing Multiple Categorical Variables

survived 0 1

pclass

1	123	200
2	158	119
3	528	181

# Visualizing Multiple Categorical Variables

survived	0	1
pclass		
1	123	200
2	158	119
3	528	181

From a crosstab, we can make a bar plot to visualize the data.

# Visualizing Multiple Categorical Variables

	survived	0	1
pclass			
1	123	200	
2	158	119	
3	528	181	

From a crosstab, we can make a bar plot to visualize the data.

```
(df[["pclass", "survived"]].value_counts().  
unstack().  
plot.bar())
```

# Visualizing Multiple Categorical Variables

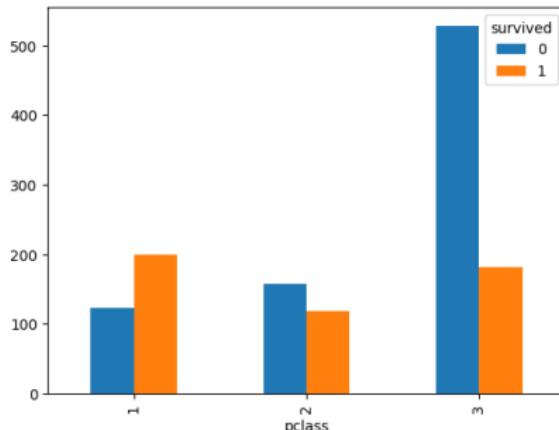
```
survived    0    1
```

```
pclass
```

	1	2	3
1	123	158	528
2	200	119	181

From a crosstab, we can make a bar plot to visualize the data.

```
(df[["pclass", "survived"]].value_counts().  
unstack().  
plot.bar())
```



# Visualizing Multiple Categorical Variables

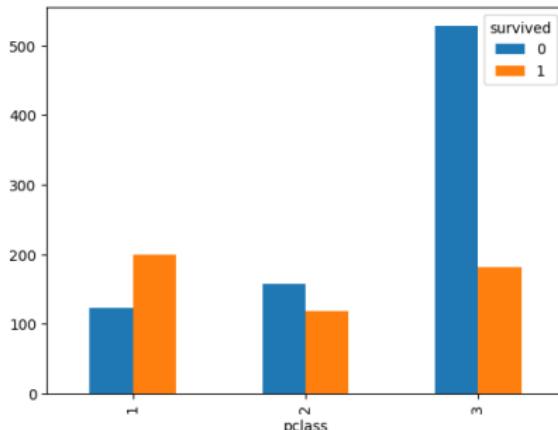
```
survived    0    1
```

```
pclass
```

1	123	200
2	158	119
3	528	181

From a crosstab, we can make a bar plot to visualize the data.

```
(df[["pclass", "survived"]].value_counts().  
unstack().  
plot.bar())
```



This is called a **grouped bar plot**.

# Marginal Counts

How do we recover the counts for each individual variable from a crosstab?

```
crosstab = df[["pclass", "survived"]].value_counts().unstack()  
crosstab
```

survived	0	1
pclass		
1	123	200
2	158	119
3	528	181

# Marginal Counts

How do we recover the counts for each individual variable from a crosstab?

```
crosstab = df[["pclass", "survived"]].value_counts().unstack()  
crosstab
```

survived	0	1
pclass		
1	123	200
2	158	119
3	528	181

We could sum over the columns (across each row) to obtain the counts for `pclass`...

```
crosstab.sum(axis="columns")
```

# Marginal Counts

How do we recover the counts for each individual variable from a crosstab?

```
crosstab = df[["pclass", "survived"]].value_counts().unstack()  
crosstab
```

	survived	0	1
pclass			
1	123	200	
2	158	119	
3	528	181	

We could sum over the columns (across each row) to obtain the counts for `pclass`...

```
crosstab.sum(axis="columns")
```

pclass	
1	323
2	277
3	709

dtype: int64

# Marginal Counts

How do we recover the counts for each individual variable from a crosstab?

```
crosstab = df[["pclass", "survived"]].value_counts().unstack()  
crosstab
```

```
survived    0    1  
  
pclass  
1      123  200  
2      158  119  
3      528  181
```

...or sum over the rows (down each column) to obtain the counts for **survived**.

```
crosstab.sum(axis="rows")
```

# Marginal Counts

How do we recover the counts for each individual variable from a crosstab?

```
crosstab = df[["pclass", "survived"]].value_counts().unstack()  
crosstab
```

```
survived    0    1  
  
pclass  
1      123  200  
2      158  119  
3      528  181
```

...or sum over the rows (down each column) to obtain the counts for **survived**.

```
crosstab.sum(axis="rows")
```

```
survived  
0      809  
1      500  
dtype: int64
```

- 1 Review
- 2 Two (or More) Categorical Variables
- 3 Proportions and Probabilities
- 4 Joint and Conditional Distributions

# Proportions

# Proportions

Instead of counts, it can be useful to report **proportions**, where we normalize by the total.

$$\text{proportion} = \frac{\text{count}}{\text{total}}.$$

# Proportions

Instead of counts, it can be useful to report **proportions**, where we normalize by the total.

$$\text{proportion} = \frac{\text{count}}{\text{total}}.$$

For example, the proportions of the three passenger classes are:

```
df["pclass"].value_counts() / len(df)
```

# Proportions

Instead of counts, it can be useful to report **proportions**, where we normalize by the total.

$$\text{proportion} = \frac{\text{count}}{\text{total}}.$$

For example, the proportions of the three passenger classes are:

```
df["pclass"].value_counts() / len(df)
```

```
3    0.541635
1    0.246753
2    0.211612
Name: pclass, dtype: float64
```

# Proportions

Instead of counts, it can be useful to report **proportions**, where we normalize by the total.

$$\text{proportion} = \frac{\text{count}}{\text{total}}.$$

For example, the proportions of the three passenger classes are:

```
df["pclass"].value_counts() / len(df)
```

```
3    0.541635
1    0.246753
2    0.211612
Name: pclass, dtype: float64
```

Together, the proportions of a categorical variable are called the **distribution** of the variable **pclass**.

# Proportions

Instead of counts, it can be useful to report **proportions**, where we normalize by the total.

$$\text{proportion} = \frac{\text{count}}{\text{total}}.$$

For example, the proportions of the three passenger classes are:

```
df["pclass"].value_counts() / len(df)
```

```
3    0.541635  
1    0.246753  
2    0.211612  
Name: pclass, dtype: float64
```

Notice that the values in a distribution add up to 1.0!

Together, the proportions of a categorical variable are called the **distribution** of the variable **pclass**.

# Probabilities

What does it mean to say, “The proportion of passengers in 3rd class is 0.541635?”

# Probabilities

What does it mean to say, “The proportion of passengers in 3rd class is 0.541635?”

One interpretation is as a **probability**.

# Probabilities

- What does it mean to say, “The proportion of passengers in 3rd class is 0.541635?”
- One interpretation is as a **probability**.
- “If we were to pick a passenger on the Titanic at random, the probability that they are in 3rd class is 0.541635.”

# Probabilities

- What does it mean to say, “The proportion of passengers in 3rd class is 0.541635?”
- One interpretation is as a **probability**.
- “If we were to pick a passenger on the Titanic at random, the probability that they are in 3rd class is 0.541635.”
- We notate this as
  - $P(\text{3rd class}) = 0.541635$

# Probabilities

- What does it mean to say, “The proportion of passengers in 3rd class is 0.541635?”
- One interpretation is as a **probability**.
- “If we were to pick a passenger on the Titanic at random, the probability that they are in 3rd class is 0.541635.”
- We note this as
  - $P(\text{3rd class}) = 0.541635$
  - or, if we want to be explicit about the variable and the category,
    - $P(\text{pclass} = 3) = 0.541635$

# Vectorization

Let's take a closer look at the code for calculating the proportions.

```
df["pclass"].value_counts() / len(df)
```

# Vectorization

Let's take a closer look at the code for calculating the proportions.

```
df["pclass"].value_counts() / len(df)
```

Notice that we divided a Series by a number! Is that even legal?

# Vectorization

Let's take a closer look at the code for calculating the proportions.

```
df["pclass"].value_counts() / len(df)
```

Notice that we divided a Series by a number! Is that even legal?

In pandas, operations are **vectorized**. A Series behaves like a vector.

# Vectorization

Let's take a closer look at the code for calculating the proportions.

```
df["pclass"].value_counts() / len(df)
```

Notice that we divided a Series by a number! Is that even legal?

In pandas, operations are **vectorized**. A Series behaves like a vector.

Vectors in pandas work like vectors in math!

# Vectorization

Let's take a closer look at the code for calculating the proportions.

```
df["pclass"].value_counts() / len(df)
```

Notice that we divided a Series by a number! Is that even legal?

In pandas, operations are **vectorized**. A Series behaves like a vector.

Vectors in pandas work like vectors in math!

## *Math Review*

To multiply a vector

$$\vec{v} = (v_1, v_2, \dots, v_n)$$

by a scalar (a.k.a. number)  $a$ ,

$$a\vec{v} =$$

# Vectorization

Let's take a closer look at the code for calculating the proportions.

```
df["pclass"].value_counts() / len(df)
```

Notice that we divided a Series by a number! Is that even legal?

In pandas, operations are **vectorized**. A Series behaves like a vector.

Vectors in pandas work like vectors in math!

## *Math Review*

To multiply a vector

$$\vec{v} = (v_1, v_2, \dots, v_n)$$

by a scalar (a.k.a. number)  $a$ ,

$$a\vec{v} = (av_1, av_2, \dots, av_n),$$

we multiply each component of  
the vector by  $a$ .

- 1 Review
- 2 Two (or More) Categorical Variables
- 3 Proportions and Probabilities
- 4 Joint and Conditional Distributions

# Joint Distributions

We can also calculate the distribution of multiple variables, called a **joint distribution**.

# Joint Distributions

We can also calculate the distribution of multiple variables, called a **joint distribution**.

```
df[["pclass", "survived"]].value_counts().unstack() / len(df)
```

# Joint Distributions

We can also calculate the distribution of multiple variables, called a **joint distribution**.

```
df[["pclass", "survived"]].value_counts().unstack() / len(df)
```

survived	0	1
pclass		
1	0.093965	0.152788
2	0.120703	0.090909
3	0.403361	0.138273

# Joint Distributions

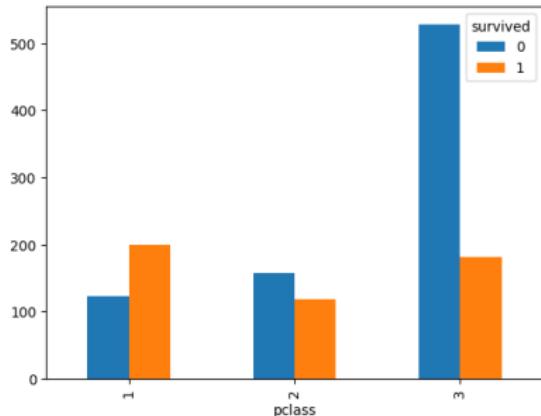
We can also calculate the distribution of multiple variables, called a **joint distribution**.

```
df[["pclass", "survived"]].value_counts().unstack() / len(df)
```

survived	0	1
pclass		
1	0.093965	0.152788
2	0.120703	0.090909
3	0.403361	0.138273

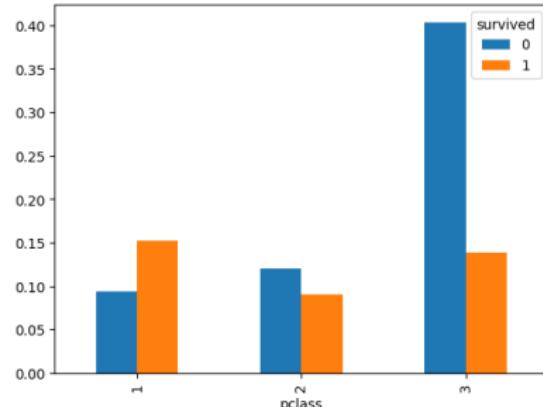
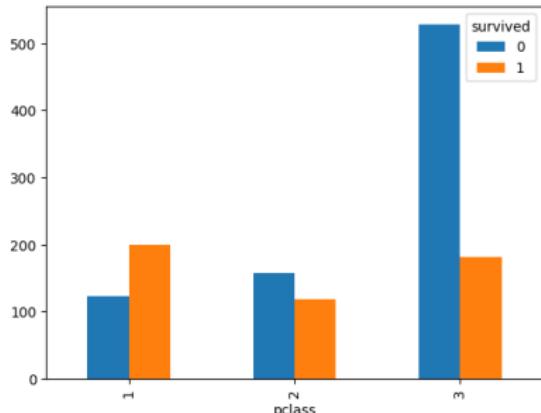
Notice that the values in the joint distribution also sum to 1.0!

# Visualizing Joint Distributions



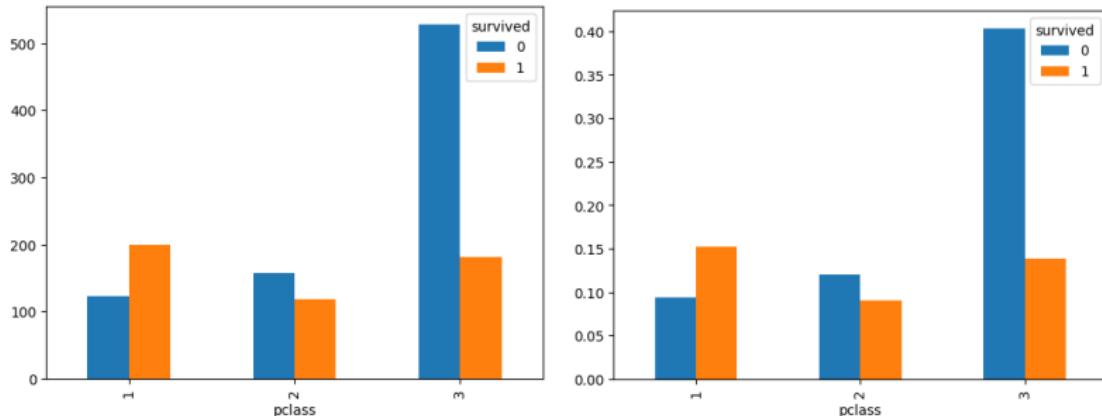
How does the bar plot change if we plot the joint distribution instead of the counts?

# Visualizing Joint Distributions



How does the bar plot change if we plot the joint distribution instead of the counts?

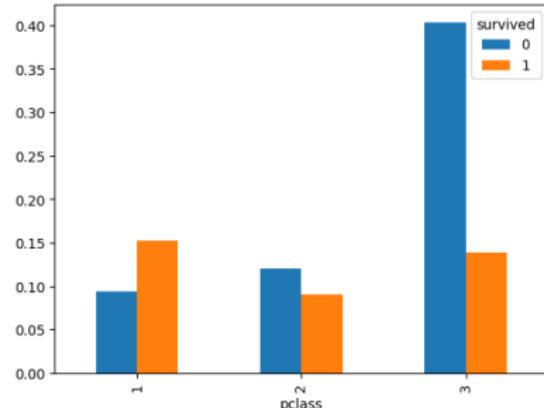
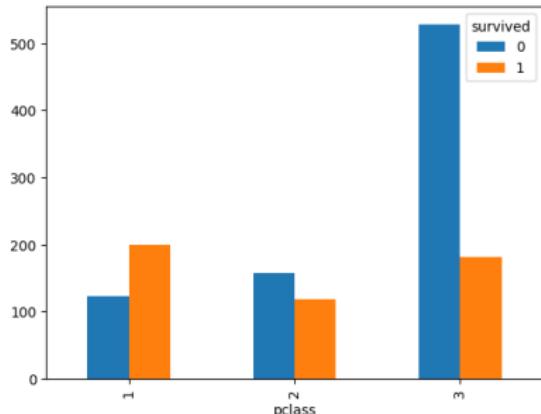
# Visualizing Joint Distributions



How does the bar plot change if we plot the joint distribution instead of the counts?

The y-axis scale changes, but the shape is the same.

# Visualizing Joint Distributions



How does the bar plot change if we plot the joint distribution instead of the counts?

The y-axis scale changes, but the shape is the same.

To appreciate the power of proportions, we need to look at conditional distributions.

# Conditional Distributions

To compare survival across the classes, we should normalize by the total in each class.

```
survived    0    1
```

```
pclass
```

	0	1
1	123	200
2	158	119
3	528	181

```
crosstab
```

# Conditional Distributions

To compare survival across the classes, we should normalize by the total in each class.

```
survived    0    1
```

survived	0	1
pclass		pclass
1	123	200
2	158	119
3	528	181

```
dtype: int64
```

```
crosstab
```

# Conditional Distributions

To compare survival across the classes, we should normalize by the total in each class.

```
survived    0    1
```

	pclass		pclass	
1	123	200	1	323
2	158	119	2	277
3	528	181	3	709

```
dtype: int64
```

```
crosstab
```

```
pclass_marginal = crosstab.sum(axis="columns")
```

# Conditional Distributions

To compare survival across the classes, we should normalize by the total in each class.

```
survived    0    1
```

pclass		
	0	1
1	123	200
2	158	119
3	528	181

pclass	
1	323
2	277
3	709

```
dtype: int64
```

```
crosstab
```

```
pclass_marginal = crosstab.sum(axis="columns")
```

Next, we divide the crosstab by the total in each class.

# Conditional Distributions

To compare survival across the classes, we should normalize by the total in each class.

```
survived    0    1  
  
pclass  
-----  
 1      123  200  
 2      158  119  
 3      528  181  
  
pclass  
-----  
 1      323  
 2      277  
 3      709  
  
dtype: int64
```

```
crosstab          pclass_marginal = crosstab.sum(axis="columns")
```

Next, we divide the crosstab by the total in each class.

```
crosstab.divide(pclass_marginal, axis="rows")
```

```
survived    0    1  
  
pclass  
-----  
 1      0.380805  0.619195  
 2      0.570397  0.429603  
 3      0.744711  0.255289
```

# Conditional Distributions

To compare survival across the classes, we should normalize by the total in each class.

```
survived    0    1  
  
pclass  
-----  
 1      123  200  
 2      158  119  
 3      528  181  
  
pclass  
-----  
 1      323  
 2      277  
 3      709  
dtype: int64
```

```
crosstab          pclass_marginal = crosstab.sum(axis="columns")
```

Next, we divide the crosstab by the total in each class.

```
crosstab.divide(pclass_marginal, axis="rows")
```

```
survived    0    1  
  
pclass  
-----  
 1      0.380805  0.619195  
 2      0.570397  0.429603  
 3      0.744711  0.255289
```

These are the **conditional distributions** of survived given pclass.

# Visualizing Conditional Distributions

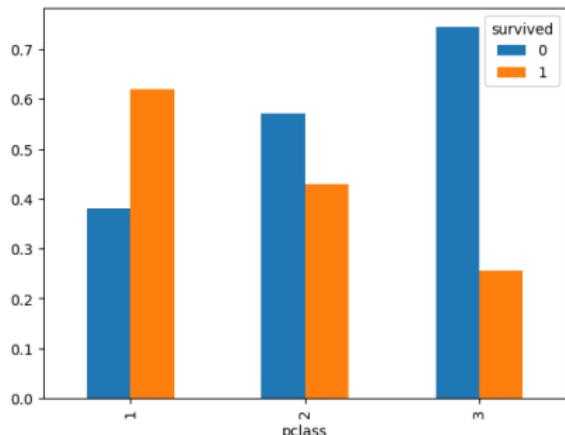
To visualize the conditional distributions, we could make a grouped bar plot...

```
(crosstab.divide(pclass_marginal, axis="rows").  
plot.bar())
```

# Visualizing Conditional Distributions

To visualize the conditional distributions, we could make a grouped bar plot...

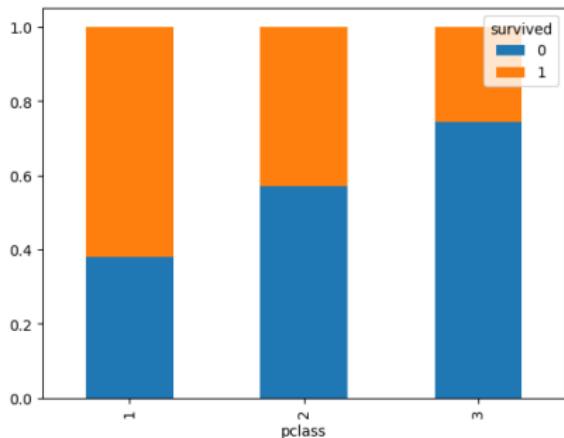
```
(crosstab.divide(pclass_marginal, axis="rows").  
    plot.bar())
```



# Visualizing Conditional Distributions

To visualize a conditional distribution, we could make a grouped bar plot...

```
(crosstab.divide(pclass_marginal, axis="rows").  
    plot.bar(stacked=True))
```



...but it is better to make a **stacked bar plot**.

# Conditional Probabilities

What does it mean to say, “The conditional proportion of survival given 3rd class is 0.255289”?

## Conditional Probabilities

What does it mean to say, “The conditional proportion of survival given 3rd class is 0.255289”?

One interpretation is as a **conditional probability**.

# Conditional Probabilities

- What does it mean to say, “The conditional proportion of survival given 3rd class is 0.255289”?
- One interpretation is as a **conditional probability**.
- “If we were to pick a 3rd class passenger on the Titanic at random, the probability that they survived is 0.255289.”

# Conditional Probabilities

- What does it mean to say, “The conditional proportion of survival given 3rd class is 0.255289”?
- One interpretation is as a **conditional probability**.
- “If we were to pick a 3rd class passenger on the Titanic at random, the probability that they survived is 0.255289.”
- We notate this as
  - $P(\text{survived}|\text{3rd class}) = 0.255289$

# Conditional Probabilities

- What does it mean to say, “The conditional proportion of survival given 3rd class is 0.255289”?
- One interpretation is as a **conditional probability**.
- “If we were to pick a 3rd class passenger on the Titanic at random, the probability that they survived is 0.255289.”
- We notate this as
  - $P(\text{survived}|\text{3rd class}) = 0.255289$