

# **Lecture 6**

## **Distances between Observations**

# Ames House Prices

## **Today's Data:** Houses in Ames, IA

Area measurements (Gr Liv Area, Lot Area, Garage Area, Total Bsmt SF) → sq ft

Count / unit (Bedroom AbvGr, Full Bath, Garage Cars) → integer / unit

Price → USD

Historical information (Year Built, Yr Sold) → numeric / year

# Ames House Prices

**Today's Data:** Houses in Ames, IA

```
import pandas as pd  
df = pd.read_table("https://datasci112.stanford.edu/data/housing.tsv")
```

# Ames House Prices

**Today's Data:** Houses in Ames, IA

Tab-Separated Values

```
import pandas as pd  
df = pd.read_table("https://datasci112.stanford.edu/data/housing.tsv")  
df
```

|      | PID       | Gr Liv Area | Bedroom | AbvGr | Full Bath | Half Bath | ... | Mo Sold | Yr Sold | Sale Type | Sale Condition | SalePrice |
|------|-----------|-------------|---------|-------|-----------|-----------|-----|---------|---------|-----------|----------------|-----------|
| 0    | 526301100 | 1656        | 3       | 1     | 0         | ...       | 5   | 2010    | WD      | Normal    | 215000         |           |
| 1    | 526350040 | 896         | 2       | 1     | 0         | ...       | 6   | 2010    | WD      | Normal    | 105000         |           |
| 2    | 526351010 | 1329        | 3       | 1     | 1         | ...       | 6   | 2010    | WD      | Normal    | 172000         |           |
| 3    | 526353030 | 2110        | 3       | 2     | 1         | ...       | 4   | 2010    | WD      | Normal    | 244000         |           |
| 4    | 527105010 | 1629        | 3       | 2     | 1         | ...       | 3   | 2010    | WD      | Normal    | 189900         |           |
| ...  | ...       | ...         | ...     | ...   | ...       | ...       | ... | ...     | ...     | ...       | ...            |           |
| 2925 | 923275080 | 1003        | 3       | 1     | 0         | ...       | 3   | 2006    | WD      | Normal    | 142500         |           |
| 2926 | 923276100 | 902         | 2       | 1     | 0         | ...       | 6   | 2006    | WD      | Normal    | 131000         |           |
| 2927 | 923400125 | 970         | 3       | 1     | 0         | ...       | 7   | 2006    | WD      | Normal    | 132000         |           |
| 2928 | 924100070 | 1389        | 2       | 1     | 0         | ...       | 4   | 2006    | WD      | Normal    | 170000         |           |
| 2929 | 924151050 | 2000        | 3       | 2     | 1         | ...       | 11  | 2006    | WD      | Normal    | 188000         |           |

2930 rows × 81 columns

Gr Liv Area: Ground living area  
Bedroom AbvGr: Bedroom above ground  
Mo Sold: Month sold  
WD: Waranty deed (normal, garantili tapu)

E.g. Gr Liv Area = 2000 → the above-ground living area is 2000 sq ft (approximately 186 m<sup>2</sup>)

## Square Footage and Bedrooms

We've seen how to visualize and summarize the relationship between two quantitative variables.

## Square Footage and Bedrooms


We've seen how to visualize and summarize the relationship between two quantitative variables.

```
df.plot.scatter(x="Gr Liv Area", y="Bedroom AbvGr", alpha=0.4)  
df[["Gr Liv Area", "Bedroom AbvGr"]].corr()
```

# Square Footage and Bedrooms

We've seen how to visualize and summarize the relationship between two quantitative variables.

```
df.plot.scatter(x="Gr Liv Area", y="Bedroom AbvGr", alpha=0.4)  
df[["Gr Liv Area", "Bedroom AbvGr"]].corr()
```



Type: Scatter plot  
X-axis: Gr Liv Area (living area in sq ft)  
Y-axis: Bedroom AbvGr (#bedrooms above ground)  
Trend: More bedrooms → generally larger living area  
Variation: Large spread within each bedroom count  
Density: Most houses have 2–4 bedrooms  
Outliers: Few very large houses with high Gr Liv Area  
every dot a house

- 1 Distances between Observations
- 2 Variable Scaling
- 3 Calculating Distances in Scikit-Learn
- 4 Summary


# Quantifying Similarity

How do we measure how “similar” two observations are?

# Quantifying Similarity


How do we measure how “similar” two observations are?

For example, how would we quantify how similar the following two houses are?




## Calculating Distance

One way is to calculate the distance between the two points.




## Calculating Distance

One way is to calculate the distance between the two points.



## Calculating Distance

One way is to calculate the distance between the two points.



By the Pythagorean Theorem,

$$d^2 = (x_1 - x'_1)^2 + (x_2 - x'_2)^2$$

or, equivalently,

$$d = \sqrt{(x_1 - x'_1)^2 + (x_2 - x'_2)^2}.$$

## Calculating Distance in Higher Dimensions

The distance formula generalizes to any number of variables  $m$ .

$$d(\mathbf{x}, \mathbf{x}') = \sqrt{\sum_{j=1}^m (x_j - x'_j)^2}$$

Distance  $\approx 0$ : Houses have almost identical features.

Moderate distance: Houses are similar but with some differences.

Very large distance: Houses differ greatly in size or room counts.

## Calculating Distance in Higher Dimensions

The distance formula generalizes to any number of variables  $m$ .

$$d(\mathbf{x}, \mathbf{x}') = \sqrt{\sum_{j=1}^m (x_j - x'_j)^2}$$


Let's calculate the distance between these two houses, based on  $m = 4$  variables: square footage, bedrooms, full baths, and half baths.



- 1 Distances between Observations
- 2 Variable Scaling
- 3 Calculating Distances in Scikit-Learn
- 4 Summary


# The Importance of Scale

This plot is misleading.




# The Importance of Scale

This plot is misleading.




The data really looks like this.




# The Importance of Scale

This plot is misleading.




The data really looks like this.




- The variation in square footage completely swamps the variation in the number of bedrooms.

# The Importance of Scale

This plot is misleading.



The data really looks like this.



The variation in square footage completely swamps the variation in the number of bedrooms.

We should bring all the variables to the same scale before calculating distance.

# Standardization

We've seen one way to bring different variables to the same scale:  
**standardization.**

Standardization is a type of scaling, similar to normalization

but it uses the mean (0) and standard deviation (1) instead of compressing values into a fixed range.

standard normal dist. ( $\text{mean}=0$ ,  $\text{sd}=1$ ) vs ( $\text{mean} \neq 0$  and  $\text{sd} \neq 1$ )

# Standardization

We've seen one way to bring different variables to the same scale:  
**standardization**.

$$x_i \leftarrow \frac{x_i - \bar{\mathbf{x}}}{\text{sd}(\mathbf{x})}$$

```
features = ["Gr Liv Area", "Bedroom AbvGr", "Full Bath", "Half Bath"]
df_standardized = ((df[features] - df[features].mean(axis="rows")) /
                    df[features].std(axis="rows"))
df_standardized
```

obtain a new DataFrame where each column has a  
**mean of 0** and a **standard deviation of 1**.

# Standardization

We've seen one way to bring different variables to the same scale:  
**standardization.**

$$x_i \leftarrow \frac{x_i - \bar{x}}{\text{sd}(\mathbf{x})}$$

```
features = ["Gr Liv Area", "Bedroom AbvGr", "Full Bath", "Half Bath"]
df_standardized = ((df[features] - df[features].mean(axis="rows")) /
                    df[features].std(axis="rows"))
df_standardized
```

|      | Gr Liv Area | Bedroom AbvGr | Full Bath | Half Bath |
|------|-------------|---------------|-----------|-----------|
| 0    | 0.309212    | 0.176064      | -1.024618 | -0.755074 |
| 1    | -1.194223   | -1.032058     | -1.024618 | -0.755074 |
| 2    | -0.337661   | 0.176064      | -1.024618 | 1.234464  |
| 3    | 1.207317    | 0.176064      | 0.783894  | 1.234464  |
| 4    | 0.255801    | 0.176064      | 0.783894  | 1.234464  |
| ...  | ...         | ...           | ...       | ...       |
| 2925 | -0.982555   | 0.176064      | -1.024618 | -0.755074 |
| 2926 | -1.182354   | -1.032058     | -1.024618 | -0.755074 |
| 2927 | -1.047836   | 0.176064      | -1.024618 | -0.755074 |
| 2928 | -0.218968   | -1.032058     | -1.024618 | -0.755074 |
| 2929 | 0.989715    | 0.176064      | 0.783894  | 1.234464  |

2930 rows × 4 columns

## The Effect of Standardization

What did standardization do?

Each feature now has mean = 0.

Each feature now has standard deviation = 1.

All features are on the same scale.

Prevents features with larger ranges from dominating analyses.

# The Effect of Standardization


What did standardization do?

```
df_standardized.plot.scatter(  
    x="Gr_Liv_Area", y="Bedroom_AbvGr", alpha=0.4)
```

# The Effect of Standardization

What did standardization do?


```
df_standardized.plot.scatter(  
    x="Gr Liv Area", y="Bedroom AbvGr", alpha=0.4)
```



# The Effect of Standardization

What did standardization do?

```
df_standardized.plot.scatter(  
    x="Gr Liv Area", y="Bedroom AbvGr", alpha=0.4)
```



Standardization ensures that every variable has mean 0 and standard deviation 1 so that every variable is treated equally.

## Finding the Most Similar Home

Let's find the most similar house by calculating distances on the *standardized* data.

# Finding the Most Similar Home

Let's find the most similar house by calculating distances on the *standardized* data.

```
import numpy as np  
diffs = df_standardized - df_standardized.loc[1707]  
dists = np.sqrt((diffs ** 2).sum(axis="columns"))  
dists.sort_values()
```

Computes the difference between each row and row 1707.

Squares the differences and sums them across columns.

Takes the square root to get Euclidean distances.

Sorts distances from closest to farthest.

# Finding the Most Similar Home

Let's find the most similar house by calculating distances on the *standardized* data.

```
import numpy as np  
diffs = df_standardized - df_standardized.loc[1707]  
dists = np.sqrt((diffs ** 2).sum(axis="columns"))  
dists.sort_values()
```

```
1707      0.000000  
160       0.043521  
909       0.249254  
2592      0.625113  
585       0.625113  
...  
2880      8.024014  
1385      8.050646  
2279      8.095655  
158       8.225392  
2723      8.313887  
Length: 2930, dtype: float64
```

## Finding the Most Similar Home

Let's find the most similar house by calculating distances on the *standardized* data.

```
import numpy as np  
diffs = df_standardized - df_standardized.loc[1707]  
dists = np.sqrt((diffs ** 2).sum(axis="columns"))  
dists.sort_values()
```

1707 0.000000

160 0.043521

909 0.249254

2592 0.625113

585 0.625113

...

2880 8.024014

1385 8.050646

2279 8.095655

158 8.225392

2723 8.313887

Length: 2930, dtype: float64

Now the most similar house to #1707 is #160, which matches our intuition when we looked at the data.

- 1 Distances between Observations
- 2 Variable Scaling
- 3 Calculating Distances in Scikit-Learn
- 4 Summary

## Scaling in Scikit-Learn

The machine learning library Scikit-Learn can be used to scale variables and calculate distances.

# Scaling in Scikit-Learn

- The machine learning library Scikit-Learn can be used to scale variables and calculate distances.

```
from sklearn.preprocessing import StandardScaler  
  
scaler = StandardScaler()  
  
# This calculates the mean and standard deviation of each variable.  
scaler.fit(df[features])  
  
# This actually does the standardization.  
array_standardized = scaler.transform(df[features])  
array_standardized
```

## Scaling in Scikit-Learn

The machine learning library Scikit-Learn can be used to scale variables and calculate distances.

```
from sklearn.preprocessing import StandardScaler  
  
scaler = StandardScaler()  
  
# This calculates the mean and standard deviation of each variable.  
scaler.fit(df[features])  
  
# This actually does the standardization.  
array_standardized = scaler.transform(df[features])  
array_standardized  
  
array([[ 0.30926506,  0.17609421, -1.02479289, -0.75520269],  
       [-1.19442705, -1.03223376, -1.02479289, -0.75520269],  
       [-0.33771825,  0.17609421, -1.02479289,  1.23467491],  
       ...,  
       [-1.04801492,  0.17609421, -1.02479289, -0.75520269],  
       [-0.21900572, -1.03223376, -1.02479289, -0.75520269],  
       [ 0.9898836 ,  0.17609421,  0.7840283 ,  1.23467491]])
```

scaled numerical array

## Scaling in Scikit-Learn

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

# This calculates the mean and standard deviation of each variable.
scaler.fit(df[features])

# This actually does the standardization.
array_standardized = scaler.transform(df[features])
array_standardized

array([[ 0.30926506,  0.17609421, -1.02479289, -0.75520269],
       [-1.19442705, -1.03223376, -1.02479289, -0.75520269],
       [-0.33771825,  0.17609421, -1.02479289,  1.23467491],
       ...,
       [-1.04801492,  0.17609421, -1.02479289, -0.75520269],
       [-0.21900572, -1.03223376, -1.02479289, -0.75520269],
       [ 0.9898836 ,  0.17609421,  0.7840283 ,  1.23467491]])
```

The result is the same as before, but it's in an array, not a DataFrame!

## Calculating Distances in Scikit-Learn

Now let's calculate the distance between house #1707 and *all* houses, using Scikit-Learn.

## Calculating Distances in Scikit-Learn

Now let's calculate the distance between house #1707 and *all* houses, using Scikit-Learn.

```
from sklearn.metrics import pairwise_distances  
  
dists = pairwise_distances(array_standardized[[1707]], array_standardized)  
dists
```

## Calculating Distances in Scikit-Learn

Now let's calculate the distance between house #1707 and *all* houses, using Scikit-Learn.

```
from sklearn.metrics import pairwise_distances  
  
dists = pairwise_distances(array_standardized[[1707]], array_standardized)  
dists  
array([[4.43704819, 6.08145351, 4.41300228, ..., 5.33963844, 5.4757908 ,  
       3.06886734]])
```

# Calculating Distances in Scikit-Learn

Now let's calculate the distance between house #1707 and *all* houses, using Scikit-Learn.

```
from sklearn.metrics import pairwise_distances

dists = pairwise_distances(array_standardized[[1707]], array_standardized)
dists
array([[4.43704819, 6.08145351, 4.41300228, ...,
       5.33963844, 5.4757908 ,
       3.06886734]])
```

We should sort these values...

```
dists[0].sort() # the sorting is done in place
dists
array([[0.          , 0.04352793, 0.24929632, ...,
       8.09703665, 8.22679607,
       8.31530563]])
```

## Calculating Distances in Scikit-Learn

Now let's calculate the distance between house #1707 and *all* houses, using Scikit-Learn.

```
from sklearn.metrics import pairwise_distances

dists = pairwise_distances(array_standardized[[1707]], array_standardized)
dists
array([[4.43704819, 6.08145351, 4.41300228, ...,
       5.33963844, 5.4757908 ,
       3.06886734]])
```

We should sort these values...

```
dists[0].sort() # the sorting is done in place
dists
array([[0.          , 0.04352793, 0.24929632, ...,
       8.09703665, 8.22679607,
       8.31530563]])
```

...but `.sort()` doesn't quite give us what we want.

## Calculating Distances in Scikit-Learn

Now let's calculate the distance between house #1707 and *all* houses, using Scikit-Learn.

```
from sklearn.metrics import pairwise_distances

dists = pairwise_distances(array_standardized[[1707]], array_standardized)
dists
array([[4.43704819, 6.08145351, 4.41300228, ...,
       5.33963844, 5.4757908 ,
       3.06886734]])
```

We should sort these values...

```
dists[0].sort() # the sorting is done in place
dists
array([[0.          , 0.04352793, 0.24929632, ...,
       8.09703665, 8.22679607,
       8.31530563]])
```

...but `.sort()` doesn't quite give us what we want.

We want to know *which* house had the smallest distance, not the value of that distance.

```
dists[0].argsort()
```

# Calculating Distances in Scikit-Learn

Now let's calculate the distance between house #1707 and *all* houses, using Scikit-Learn.

```
from sklearn.metrics import pairwise_distances  
  
dists = pairwise_distances(array_standardized[[1707]], array_standardized)  
dists  
array([[4.43704819, 6.08145351, 4.41300228, ..., 5.33963844, 5.4757908 ,  
       3.06886734]])
```

argsort() returns the indices that would sort an array or Series.  
The smallest values' indices come first, largest last.  
Useful for ranking or finding nearest/farthest items without changing the original array.

We should sort these values...

```
dists[0].sort() # the sorting is done in place  
dists  
array([[0.          , 0.04352793, 0.24929632, ..., 8.09703665, 8.22679607,  
       8.31530563]])
```

...but .sort() doesn't quite give us what we want.

We want to know *which* house had the smallest distance, not the value of that distance.

```
dists[0].argsort()  
array([1707, 160, 909, ..., 2279, 158, 2723])
```

## The Distance Matrix

We can also calculate the distance between every pair of observations and store the result in a matrix.

## The Distance Matrix

We can also calculate the distance between every pair of observations and store the result in a matrix.

```
dist_matrix = pairwise_distances(array_standardized)  
dist_matrix
```

# The Distance Matrix

We can also calculate the distance between every pair of observations and store the result in a matrix.

```
dist_matrix = pairwise_distances(array_standardized)
dist_matrix
array([[0.          , 1.92902733, 2.09241494, ... , 1.35727998, 1.31875946, 2.77393016],
       [1.92902733, 0.          , 2.48064897, ... , 1.21716597, 0.97542133, 3.66915745],
       [2.09241494, 2.48064897, 0.          , ... , 2.11284979, 2.33104312, 2.24373812],
       ... ,
       [1.35727998, 1.21716597, 2.11284979, ... , 0.          , 1.4653712 , 3.37408911],
       [1.31875946, 0.97542133, 2.33104312, ... , 1.4653712 , 0.          , 3.1863642 ],
       [2.77393016, 3.66915745, 2.24373812, ... , 3.37408911, 3.1863642 , 0.        ]])
```

# The Distance Matrix

We can also calculate the distance between every pair of observations and store the result in a matrix.

```
dist_matrix = pairwise_distances(array_standardized)
dist_matrix
```

```
array([[0.          , 1.92902733, 2.09241494, ... , 1.35727998, 1.31875946, 2.77393016],
       [1.92902733, 0.          , 2.48064897, ... , 1.21716597, 0.97542133, 3.66915745],
       [2.09241494, 2.48064897, 0.          , ... , 2.11284979, 2.33104312, 2.24373812],
       ... ,
       [1.35727998, 1.21716597, 2.11284979, ... , 0.          , 1.4653712 , 3.37408911],
       [1.31875946, 0.97542133, 2.33104312, ... , 1.4653712 , 0.          , 3.1863642 ],
       [2.77393016, 3.66915745, 2.24373812, ... , 3.37408911, 3.1863642 , 0.        ]])
```

Note that row #1707 of this matrix contains the dists we calculated previously.

```
dist_matrix[1707]
```

```
array([4.43704819, 6.08145351, 4.41300228, ... , 5.33963844, 5.4757908 ,
       3.06886734])
```

- 1 Distances between Observations
- 2 Variable Scaling
- 3 Calculating Distances in Scikit-Learn
- 4 Summary

## Summary

To measure similarity between two observations, we calculate the distance between them.

## Summary

To measure similarity between two observations, we calculate the distance between them.

We need to first make sure that all the variables are on the same scale.

## Summary

To measure similarity between two observations, we calculate the distance between them.

We need to first make sure that all the variables are on the same scale.

There are actually different choices of scalings and distance metrics, which you will explore in section tomorrow.

### Scaling Variables

- Standardization

$$x_i \leftarrow \frac{x_i - \bar{\mathbf{x}}}{\text{sd}(\mathbf{x})}$$

- Min-Max Scaling

$$x_i \leftarrow \frac{x_i - \min(\mathbf{x})}{\max(\mathbf{x}) - \min(\mathbf{x})}$$

### Distance Metrics

- Euclidean ( $\ell_2$ )

$$\sqrt{\sum_{j=1}^m (x_j - x'_j)^2}$$

- Manhattan ( $\ell_1$ )

$$\sum_{j=1}^m |x_j - x'_j|$$