

# **Lecture 15**

## **Evaluating Classification Models**

## Case Study: Credit Card Fraud

Data set of credit card transactions from Vesta.

```
import pandas as pd  
  
df_fraud = pd.read_csv("https://datasci112.stanford.edu/data/fraud.csv")  
df_fraud
```

# Case Study: Credit Card Fraud

Data set of credit card transactions from Vesta.

```
import pandas as pd  
  
df_fraud = pd.read_csv("https://datasci112.stanford.edu/data/fraud.csv")  
df_fraud
```

	c1-c14: dolandırıcılık ile ilgili oluşturulan istatistiksel veya davranışsal özellikler														sahte mi değil mi				
	card4	card6	P_emaildomain	TransactionAmt	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10	C11	C12	C13	C14	isFraud
0	visa	debit	gmail.com	62.950	139.0	110.0	0.0	0.0	135.0	93.0	0.0	0.0	93.0	0.0	93.0	0.0	637.0	114.0	0
1	visa	debit		35.950	1.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0	3.0	1.0	0
2	visa	debit	yahoo.com	117.000	1.0	1.0	0.0	0.0	0.0	2.0	0.0	0.0	1.0	0.0	1.0	0.0	4.0	1.0	1
3	visa	debit	hotmail.com	54.500	1.0	1.0	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	0
4	visa	debit	gmail.com	255.000	1.0	3.0	0.0	0.0	0.0	4.0	0.0	0.0	2.0	0.0	3.0	1.0	20.0	1.0	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	
59049	mastercard	debit	gmail.com	20.522	1.0	1.0	0.0	1.0	0.0	1.0	1.0	1.0	0.0	1.0	1.0	1.0	1.0	1.0	0
59050	mastercard	credit	yahoo.com	50.000	1.0	1.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	1.0	0
59051	mastercard	debit	icloud.com	97.950	2.0	2.0	0.0	0.0	1.0	2.0	0.0	0.0	1.0	0.0	2.0	0.0	15.0	2.0	0
59052	visa	debit	gmail.com	16.723	1.0	2.0	0.0	1.0	0.0	1.0	1.0	2.0	0.0	1.0	1.0	1.0	1.0	1.0	1
59053	mastercard	debit		107.950	24.0	21.0	0.0	0.0	14.0	13.0	0.0	0.0	17.0	0.0	18.0	0.0	84.0	17.0	0

59054 rows x 19 columns

# Case Study: Credit Card Fraud

Data set of credit card transactions from Vesta.

```
import pandas as pd  
  
df_fraud = pd.read_csv("https://datasci112.stanford.edu/data/fraud.csv")  
df_fraud
```

	card4	card6	P_emaildomain	TransactionAmt	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10	C11	C12	C13	C14	isFraud
0	visa	debit	gmail.com	62.950	139.0	110.0	0.0	0.0	135.0	93.0	0.0	0.0	93.0	0.0	93.0	0.0	637.0	114.0	0
1	visa	debit		35.950	1.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0	3.0	1.0	0
2	visa	debit	yahoo.com	117.000	1.0	1.0	0.0	0.0	0.0	2.0	0.0	0.0	1.0	0.0	1.0	0.0	4.0	1.0	1
3	visa	debit	hotmail.com	54.500	1.0	1.0	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	0.0	0.0	0.0	0
4	visa	debit	gmail.com	255.000	1.0	3.0	0.0	0.0	0.0	4.0	0.0	0.0	2.0	0.0	3.0	1.0	20.0	1.0	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	
59049	mastercard	debit	gmail.com	20.522	1.0	1.0	0.0	1.0	0.0	1.0	1.0	1.0	0.0	1.0	1.0	1.0	1.0	1.0	0
59050	mastercard	credit	yahoo.com	50.000	1.0	1.0	0.0	0.0	0.0	1.0	0.0	1.0	0.0	0.0	1.0	0.0	1.0	1.0	0
59051	mastercard	debit	icloud.com	97.950	2.0	2.0	0.0	0.0	1.0	2.0	0.0	0.0	1.0	0.0	2.0	0.0	15.0	2.0	0
59052	visa	debit	gmail.com	16.723	1.0	2.0	0.0	1.0	0.0	1.0	1.0	2.0	0.0	1.0	1.0	1.0	1.0	1.0	1
59053	mastercard	debit		107.950	24.0	21.0	0.0	0.0	14.0	13.0	0.0	0.0	17.0	0.0	18.0	0.0	84.0	17.0	0

59054 rows × 19 columns

Goal: Predict isFraud, where 1 indicates a fraudulent transaction.

## ① Recap

## ② Precision and Recall

# Classification Model

We can use  $k$ -nearest neighbors for classification:

```
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.pipeline import make_pipeline
from sklearn.compose import make_column_transformer

pipeline = make_pipeline(
    make_column_transformer(
        (OneHotEncoder(handle_unknown="ignore", sparse_output=False),
         ["card4", "card6", "P_emaildomain"]),
        remainder="passthrough"),
    StandardScaler(),
    KNeighborsClassifier(n_neighbors=5))
```

"Kategorik sütunları one-hot encode et, tüm sayıları ölçekle(scale),  
sonra 5-komşulu KNN ile fraud / normal sınıflandır."

## Training a Classifier

```
X_train = df_fraud.drop("isFraud", axis="columns")  
y_train = df_fraud["isFraud"]
```

# Training a Classifier

```
X_train = df_fraud.drop("isFraud", axis="columns")
y_train = df_fraud["isFraud"]
```

```
from sklearn.model_selection import cross_val_score

cross_val_score(
    pipeline,
    X=X_train, y=y_train,
    scoring="accuracy",
    cv=10
).mean()
```

**cv=10**

**the most commonly used value** (the standard in the literature).

10-fold cross-validation reduces overfitting while keeping the computational cost reasonable.

(Bir öğrenci sınav sorularının cevaplarını ezberler → O sınavda tam puan, ama yeni sorular gelince başarısız olur.

Modelde de aynısı olur: ezberleme = overfitting.)

# Training a Classifier

```
X_train = df_fraud.drop("isFraud", axis="columns")
y_train = df_fraud["isFraud"]
```

```
from sklearn.model_selection import cross_val_score

cross_val_score(
    pipeline,
    X=X_train, y=y_train,
    scoring="accuracy",
    cv=10
).mean()
```

0.9681647131621484

# Training a Classifier

```
X_train = df_fraud.drop("isFraud", axis="columns")
y_train = df_fraud["isFraud"]
```

```
from sklearn.model_selection import cross_val_score
cross_val_score(
    pipeline,
    X=X_train, y=y_train,
    scoring="accuracy",
    cv=10
).mean()
```

0.9681647131621484

How is the accuracy so high?

## **A Closer Look**

Let's take a closer look at the labels.

## A Closer Look

Let's take a closer look at the labels.

```
y_train.value_counts()
```

## A Closer Look

Let's take a closer look at the labels.

```
y_train.value_counts()  
0    56935  
1    2119  
Name: isFraud, dtype: int64
```

## A Closer Look

Let's take a closer look at the labels.

```
y_train.value_counts()
```

```
0    56935  
1    2119  
Name: isFraud, dtype: int64
```

The vast majority of transactions are normal!

## A Closer Look

Let's take a closer look at the labels.

```
y_train.value_counts()
```

```
0    56935  
1    2119  
Name: isFraud, dtype: int64
```

The vast majority of transactions are normal!

If we just predicted that every transaction is normal, the accuracy would be

## A Closer Look

Let's take a closer look at the labels.

```
y_train.value_counts()
```

```
0    56935  
1    2119  
Name: isFraud, dtype: int64
```

The vast majority of transactions are normal!

If we just predicted that every transaction is normal, the accuracy would be  $1 - \frac{2119}{59054} = .964$ .

## A Closer Look

Let's take a closer look at the labels.

```
y_train.value_counts()
```

```
0    56935  
1    2119  
Name: isFraud, dtype: int64
```

The vast majority of transactions are normal!

If we just predicted that every transaction is normal, the accuracy would be  $1 - \frac{2119}{59054} = .964$ .

Even though such predictions would be accurate *overall*, it is inaccurate for fraudulent transactions. A good model is “accurate for every class”.

1 Recap

2 Precision and Recall

## Precision and Recall

We need a score that measures “accuracy for class  $c$ ”.

## Precision and Recall

We need a score that measures “accuracy for class  $c$ ”.

There are at least two reasonable definitions:

## Precision and Recall

We need a score that measures “accuracy for class  $c$ ”.

There are at least two reasonable definitions:

- **precision:**  $P(\text{correct}|\text{predicted class } C)$

## Precision and Recall

- We need a score that measures “accuracy for class  $C$ ”.
- There are at least two reasonable definitions:
  - **precision:**  $P(\text{correct}|\text{predicted class } C)$ 
    - Among the observations that were predicted to be in class  $C$ , what proportion actually were? Pozitif dediğinin ne kadarının gerçekten pozitif olduğu.

## Precision and Recall

- We need a score that measures “accuracy for class  $C$ ”.
- There are at least two reasonable definitions:
  - **precision:**  $P(\text{correct}|\text{predicted class } C)$ 
    - Among the observations that were predicted to be in class  $C$ , what proportion actually were?
  - **recall:**  $P(\text{correct}|\text{actual class } C)$ .

## Precision and Recall

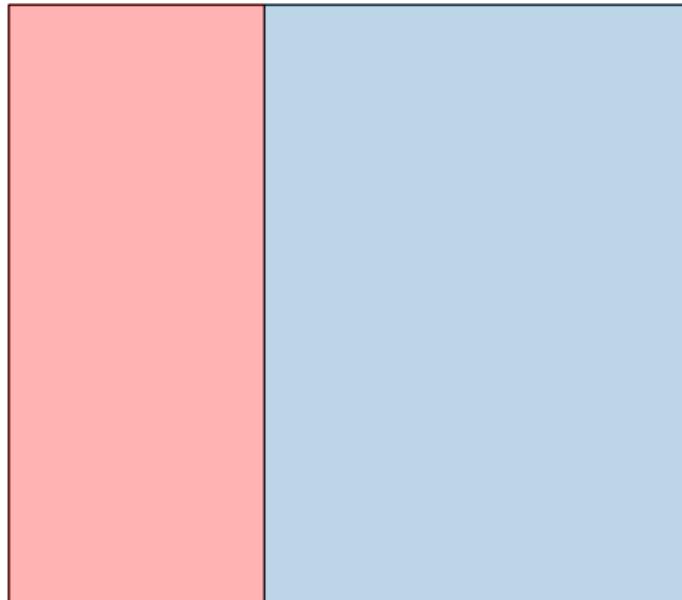
- We need a score that measures “accuracy for class  $C$ ”.
- There are at least two reasonable definitions:
  - **precision:**  $P(\text{correct}|\text{predicted class } C)$ 
    - Among the observations that were predicted to be in class  $C$ , what proportion actually were? *Pozitif dediğinin ne kadarının gerçekten pozitif olduğu.*
  - **recall:**  $P(\text{correct}|\text{actual class } C).$ 
    - Among the observations that were actually in class  $C$ , what proportion were predicted to be? *Gerçek pozitiflerin ne kadarını yakalayabildin?*

# A Geometric Look at Precision and Recall

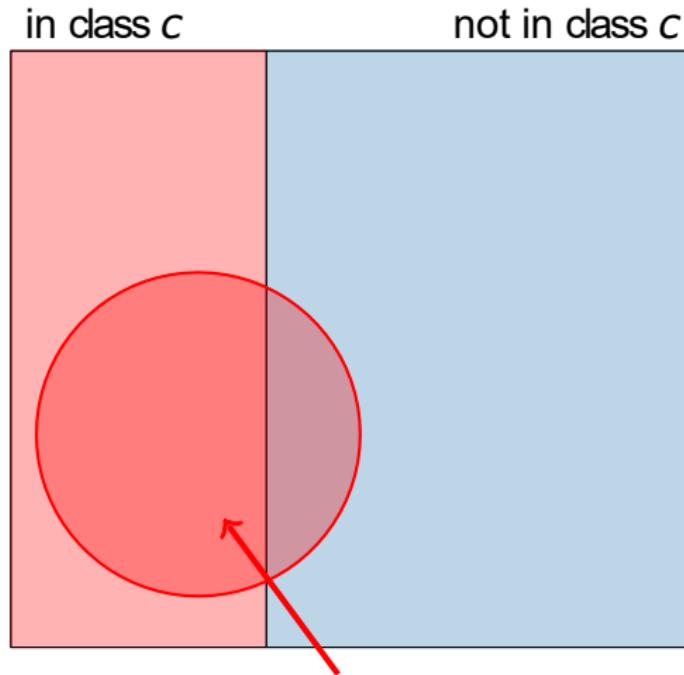
## A Geometric Look at Precision and Recall

in class  $c$

not in class  $c$

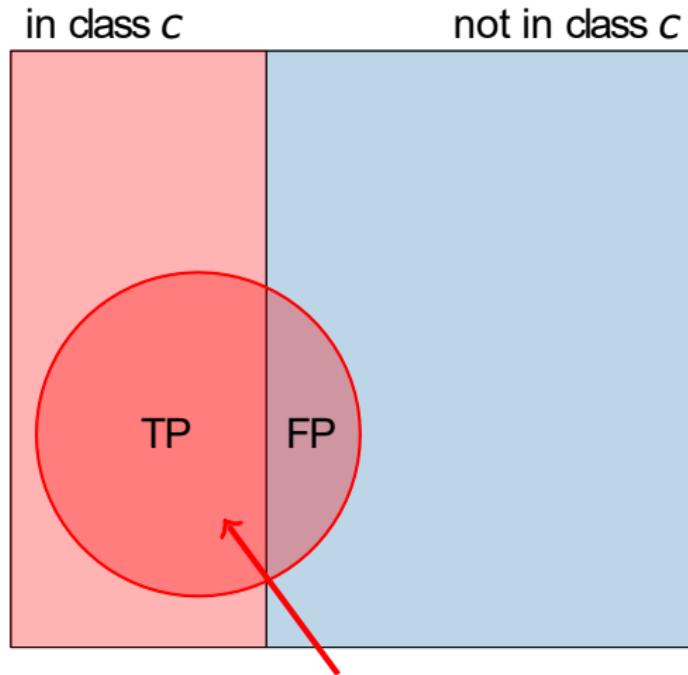


## A Geometric Look at Precision and Recall



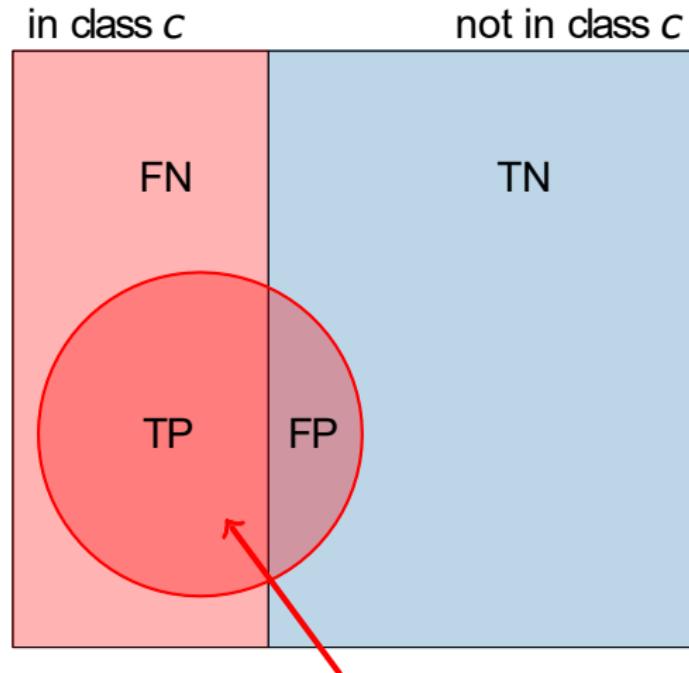
predicted to be in class c  
(a.k.a. “positives”)

## A Geometric Look at Precision and Recall



predicted to be in class c  
(a.k.a. “positives”)

## A Geometric Look at Precision and Recall

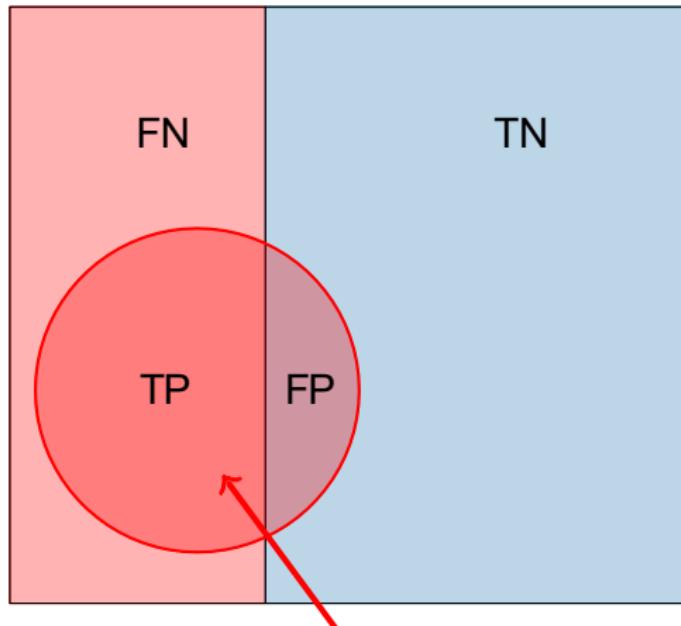


predicted to be in class c  
(a.k.a. “positives”)

# A Geometric Look at Precision and Recall

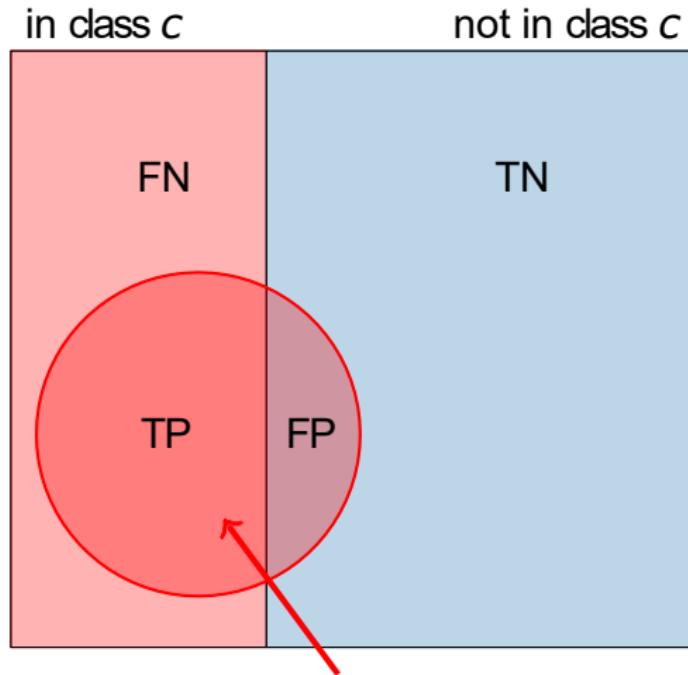
in class  $c$

not in class  $c$



**precision** = \_\_\_\_\_

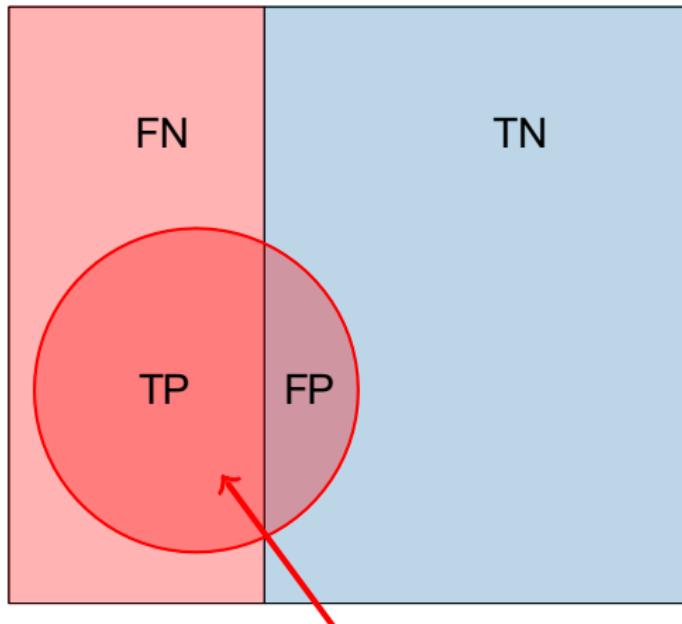
## A Geometric Look at Precision and Recall



$$\text{precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

## A Geometric Look at Precision and Recall

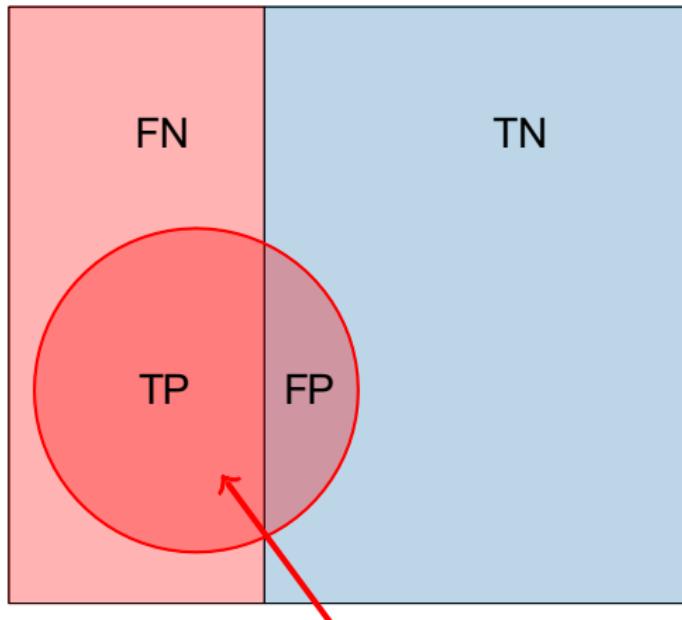
in class  $c$       not in class  $c$



$$\text{precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

## A Geometric Look at Precision and Recall

in class  $c$       not in class  $c$



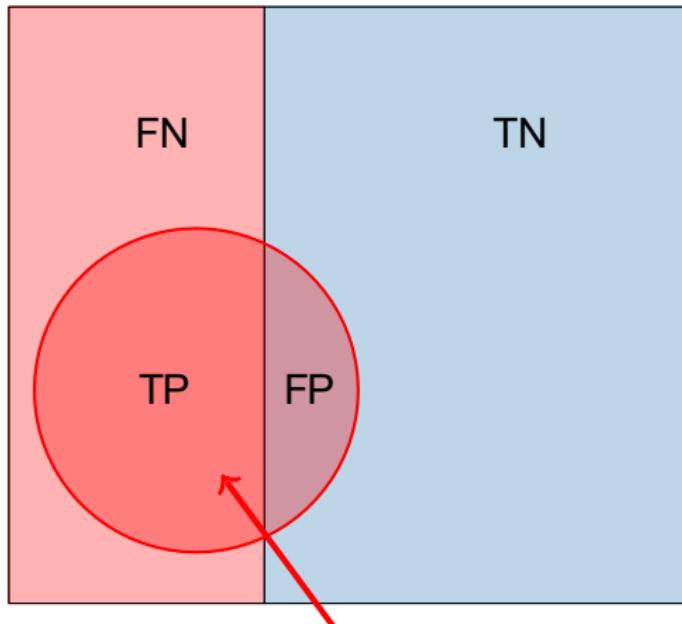
predicted to be in class  $c$   
(a.k.a. "positives")

$$\text{precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$
       $\text{recall} = \underline{\hspace{2cm}}$

# A Geometric Look at Precision and Recall

in class  $c$

not in class  $c$



predicted to be in class  $c$   
(a.k.a. “positives”)

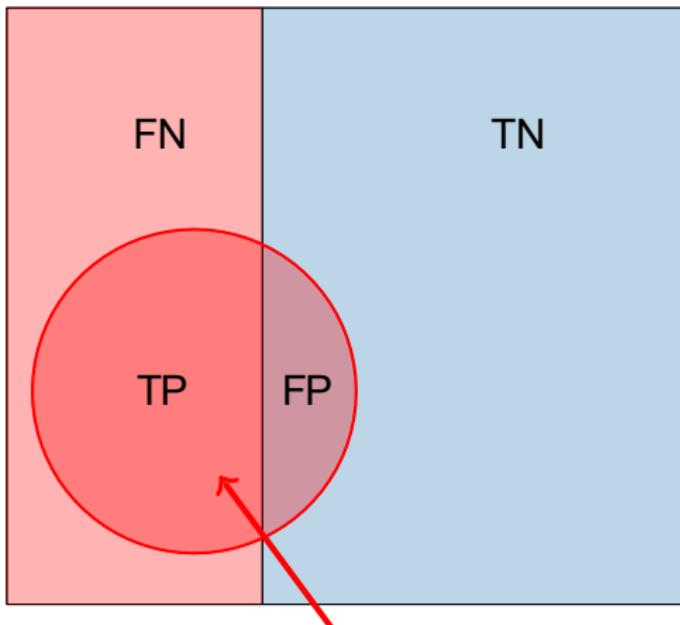
$$\text{precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

$$\text{recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

# A Geometric Look at Precision and Recall

in class  $c$

not in class  $c$



predicted to be in class  $c$   
(a.k.a. “positives”)

Ne kadar isabetli pozitif tahmin yaptım?

$$\text{precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

$$\text{recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

Gerçek pozitiflerin ne kadarını yakalayabildim?

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

```
from sklearn.metrics import confusion_matrix
pipeline.fit(X_train, y_train)
y_train_ = pipeline.predict(X_train)
confusion_matrix(y_train, y_train_)
```

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

```
from sklearn.metrics import confusion_matrix
pipeline.fit(X_train, y_train)
y_train_ = pipeline.predict(X_train)
confusion_matrix(y_train, y_train_)
```

```
array([[56817,    118],
       [ 1524,    595]])
```

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

```
from sklearn.metrics import confusion_matrix
pipeline.fit(X_train, y_train)
y_train_ = pipeline.predict(X_train)
confusion_matrix(y_train, y_train_)
```

```
array([[56817,    118],      ← actually in class 0
       [ 1524,    595]])
```

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

```
from sklearn.metrics import confusion_matrix
pipeline.fit(X_train, y_train)
y_train_ = pipeline.predict(X_train)
confusion_matrix(y_train, y_train_)
```

```
array([[56817,    118],      ← actually in class 0
       [ 1524,    595]])    ← actually in class 1
```

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

```
from sklearn.metrics import confusion_matrix
pipeline.fit(X_train, y_train)
y_train_ = pipeline.predict(X_train)
confusion_matrix(y_train, y_train_)
```

array([[56817, 118],	← actually in class 0
[ 1524, 595]])	← actually in class 1

↑  
predicted  
to be in  
class 0

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

```
from sklearn.metrics import confusion_matrix
pipeline.fit(X_train, y_train)
y_train_ = pipeline.predict(X_train)
confusion_matrix(y_train, y_train_)
```

array([[56817, 118],  
 [ 1524, 595]])  
 ↑ ↑  
predicted predicted  
to be in to be in  
class 0 class 1

← actually in class 0  
← actually in class 1

	Predicted 0	Predicted 1
Actual 0	56817	118
Actual 1	1524	595

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

```
from sklearn.metrics import confusion_matrix
pipeline.fit(X_train, y_train)
y_train_ = pipeline.predict(X_train)
confusion_matrix(y_train, y_train_)
```

	Predicted 0	Predicted 1
Actual 0	TN=56817	FP=118
Actual 1	FN=1524	TP=595

```
array([[56817,    118],      ← actually in class 0
       [ 1524,    595]])    ← actually in class 1
```

- What is the (training) accuracy?

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

```
from sklearn.metrics import confusion_matrix
pipeline.fit(X_train, y_train)
y_train_ = pipeline.predict(X_train)
confusion_matrix(y_train, y_train_)
```

	Predicted 0	Predicted 1
Actual 0	TN=56817	FP=118
Actual 1	FN=1524	TP=595

```
array([[56817,    118],      ← actually in class 0
       [ 1524,    595]])    ← actually in class 1
```

- What is the (training) accuracy? 97.2%

correct predictions/all

$$\text{Accuracy} = (\text{TP} + \text{TN}) / \text{Total samples}$$

$$\text{TN} = 56,817$$

$$\text{TP} = 595$$

$$\text{Total samples} = 56,817 + 1,524 + 118 + 595 = 59,054$$

$$\text{Accuracy} = (56,817 + 595) / 59,054 \approx 0.972 = 97.2\%$$

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

```
from sklearn.metrics import confusion_matrix
pipeline.fit(X_train, y_train)
y_train_ = pipeline.predict(X_train)
confusion_matrix(y_train, y_train_)
```

	Predicted 0	Predicted 1
Actual 0	TN=56817	FP=118
Actual 1	FN=1524	TP=595

```
array([[56817,    118],    ← actually in class 0
       [ 1524,   595]])    ← actually in class 1
```

- What is the (training) accuracy? 97.2%
- What's the precision for normal transactions?

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

```
from sklearn.metrics import confusion_matrix
pipeline.fit(X_train, y_train)
y_train_ = pipeline.predict(X_train)
confusion_matrix(y_train, y_train_)
```

	Predicted 0	Predicted 1
Actual 0	TN=56817	FP=118
Actual 1	FN=1524	TP=595

```
array([[56817,    118],      ← actually in class 0
       [ 1524,   595]])      ← actually in class 1
```

- What is the (training) accuracy? 97.2%
- What's the precision for normal transactions? 97.4%

TN=56817

FN=1524

$$\text{Precision} = 56817 / (56817 + 1524)$$

$$\text{Precision} = 56817 / 58341$$

$$= 0.973877 \approx 0.974 \text{ (97.4%)}$$

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

```
from sklearn.metrics import confusion_matrix
pipeline.fit(X_train, y_train)
y_train_ = pipeline.predict(X_train)
confusion_matrix(y_train, y_train_)
```

	Predicted 0	Predicted 1
Actual 0	TN=56817	FP=118
Actual 1	FN=1524	TP=595

```
array([[56817,    118],    ← actually in class 0
       [ 1524,   595]])    ← actually in class 1
```

- What is the (training) accuracy? 97.2%
- What's the precision for normal transactions? 97.4%
- What's the recall for normal transactions?

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

```
from sklearn.metrics import confusion_matrix
pipeline.fit(X_train, y_train)
y_train_ = pipeline.predict(X_train)
confusion_matrix(y_train, y_train_)
```

	Predicted 0	Predicted 1
Actual 0	TN=56817	FP=118
Actual 1	FN=1524	TP=595

```
array([[56817,    118],      ← actually in class 0
       [ 1524,    595]])      ← actually in class 1
```

- What is the (training) accuracy? 97.2%
- What's the precision for normal transactions? 97.4%
- What's the recall for normal transactions? 99.8%

```
TN=56817
FP=118
Recall_0 =56817/ (56817+118)
≈0.9979(99.8%)
```

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

```
from sklearn.metrics import confusion_matrix
pipeline.fit(X_train, y_train)
y_train_ = pipeline.predict(X_train)
confusion_matrix(y_train, y_train_)
```

	Predicted 0	Predicted 1
Actual 0	TN=56817	FP=118
Actual 1	FN=1524	TP=595

```
array([[56817,    118],    ← actually in class 0
       [ 1524,   595]])    ← actually in class 1
```

- What is the (training) accuracy? 97.2%
- What's the precision for normal transactions? 97.4%
- What's the recall for normal transactions? 99.8%
- What's the precision for fraudulent transactions?

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

```
from sklearn.metrics import confusion_matrix
pipeline.fit(X_train, y_train)
y_train_ = pipeline.predict(X_train)
confusion_matrix(y_train, y_train_)
```

	Predicted 0	Predicted 1
Actual 0	TN=56817	FP=118
Actual 1	FN=1524	TP=595

```
array([[56817,    118],    ← actually in class 0
       [ 1524,   595]])    ← actually in class 1
```

- What is the (training) accuracy? 97.2%
- What's the precision for normal transactions? 97.4%
- What's the recall for normal transactions? 99.8%
- What's the precision for fraudulent transactions? 83.5%

$$\text{Precision} = 595 / (595 + 118)$$

(actual1 and predicted1 / predicted1)

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

```
from sklearn.metrics import confusion_matrix
pipeline.fit(X_train, y_train)
y_train_ = pipeline.predict(X_train)
confusion_matrix(y_train, y_train_)
```

	Predicted 0	Predicted 1
Actual 0	TN=56817	FP=118
Actual 1	FN=1524	TP=595

```
array([[56817,    118],      ← actually in class 0
       [ 1524,   595]])      ← actually in class 1
```

- What is the (training) accuracy? 97.2%
- What's the precision for normal transactions? 97.4%
- What's the recall for normal transactions? 99.8%
- What's the precision for fraudulent transactions? 83.5%
- What's the recall for fraudulent transactions?

actual1 and predicted1/actual1

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

```
from sklearn.metrics import confusion_matrix
pipeline.fit(X_train, y_train)
y_train_ = pipeline.predict(X_train)
confusion_matrix(y_train, y_train_)
```

	Predicted 0	Predicted 1
Actual 0	TN=56817	FP=118
Actual 1	FN=1524	TP=595

array([[56817, 118],  
 [1524, 595]]) ← actually in class 0  
 ← actually in class 1

- What is the (training) accuracy? 97.2% actual true predicted/all
- What's the precision for normal transactions? 97.4% pre0andAct0/pre0
- What's the recall for normal transactions? 99.8% Act0andPre0/act0
- What's the precision for fraudulent transactions? 83.5%
- What's the recall for fraudulent transactions? 28.1%

$$\text{recall} = 595 / (595 + 1524)$$

## Precision and Recall by Hand

To check our understanding of these definitions, let's calculate a few precisions and recalls by hand.

First, summarize the results by the **confusion matrix**.

```
from sklearn.metrics import confusion_matrix
pipeline.fit(X_train, y_train)
y_train_ = pipeline.predict(X_train)
confusion_matrix(y_train, y_train_)
```

array([[56817, 118],  
 [ 1524, 595]])  
 ← actually in class 0  
 ← actually in class 1

- What is the (training) accuracy? 97.2%
- What's the precision for normal transactions? 97.4%
- What's the recall for normal transactions? 99.8%
- What's the precision for fraudulent transactions? 83.5%
- What's the recall for fraudulent transactions? 28.1%

*Note:* Each class has its own precision and recall!

## **Tradeoff between Precision and Recall**

Can you imagine a classifier that always has 100% recall for class c, no matter the data?

## Tradeoff between Precision and Recall

Can you imagine a classifier that always has 100% recall for class  $c$ , no matter the data?

In general, if the model classifies more observations as  $c$ ,

- recall (for class  $c$ )  $\uparrow$

## Tradeoff between Precision and Recall

Can you imagine a classifier that always has 100% recall for class  $c$ , no matter the data?

In general, if the model classifies more observations as  $c$ ,

- recall (for class  $c$ )  $\uparrow$
- precision (for class  $c$ )  $\downarrow$

## Tradeoff between Precision and Recall

Can you imagine a classifier that always has 100% recall for class  $c$ , no matter the data?

In general, if the model classifies more observations as  $c$ ,

- recall (for class  $c$ )  $\uparrow$
- precision (for class  $c$ )  $\downarrow$

How do we compare two classifiers, if one has higher precision and the other has higher recall?

## Tradeoff between Precision and Recall

Can you imagine a classifier that always has 100% recall for class  $c$ , no matter the data?

In general, if the model classifies more observations as  $c$ ,

- recall (for class  $c$ )  $\uparrow$
- precision (for class  $c$ )  $\downarrow$

How do we compare two classifiers, if one has higher precision and the other has higher recall?

The **F1 score** combines precision and recall into a single score:

F1 score = harmonic mean of precision and recall

## Tradeoff between Precision and Recall

Can you imagine a classifier that always has 100% recall for class  $c$ , no matter the data?

In general, if the model classifies more observations as  $c$ ,

- recall (for class  $c$ )  $\uparrow$
- precision (for class  $c$ )  $\downarrow$

How do we compare two classifiers, if one has higher precision and the other has higher recall?

The **F1 score** combines precision and recall into a single score:

F1 score = harmonic mean of precision and recall

$$= 1 / \frac{1}{2} \left( \frac{1}{\text{precision}} + \frac{1}{\text{recall}} \right)$$

## Tradeoff between Precision and Recall

Can you imagine a classifier that always has 100% recall for class  $c$ , no matter the data?

In general, if the model classifies more observations as  $c$ ,

- recall (for class  $c$ )  $\uparrow$
- precision (for class  $c$ )  $\downarrow$

How do we compare two classifiers, if one has higher precision and the other has higher recall?

The **F1 score** combines precision and recall into a single score:

$$\begin{aligned}\text{F1 score} &= \text{harmonic mean of precision and recall} \\ &= 1 / \frac{1}{2} \left( \frac{1}{\text{precision}} + \frac{1}{\text{recall}} \right)\end{aligned}$$

So the F1 score of the classifier for fraudulent transactions is

$$1 / \frac{1}{2} \left( \frac{1}{.835} + \frac{1}{.281} \right) \approx 420\%.$$

## Tradeoff between Precision and Recall

Can you imagine a classifier that always has 100% recall for class  $c$ , no matter the data?

In general, if the model classifies more observations as  $c$ ,

- recall (for class  $c$ )  $\uparrow$
- precision (for class  $c$ )  $\downarrow$

How do we compare two classifiers, if one has higher precision and the other has higher recall?

The **F1 score** combines precision and recall into a single score:

$$\begin{aligned}\text{F1 score} &= \text{harmonic mean of precision and recall} \\ &= 1 / \frac{1}{2} \left( \frac{1}{\text{precision}} + \frac{1}{\text{recall}} \right)\end{aligned}$$

So the F1 score of the classifier for fraudulent transactions is

$$1 / \frac{1}{2} \left( \frac{1}{.835} + \frac{1}{.281} \right) \approx 420\%.$$

To achieve a high F1 score, both precision and recall have to be high. If either is low, then the harmonic mean will be low.

## **Estimating Test Precision, Recall, and F1**

Remember that each class has its own precision, recall, and F1.

## **Estimating Test Precision, Recall, and F1**

Remember that each class has its own precision, recall, and F1.

But Scikit-Learn requires that the `scoring=` parameter be a single number.

## Estimating Test Precision, Recall, and F1

Remember that each class has its own precision, recall, and F1.

But Scikit-Learn requires that the `scoring=` parameter be a single number.

For this, we can use

- "precision\_macro"
- "recall\_macro"
- "f1\_macro"

which average the score over the classes.

# Estimating Test Precision, Recall, and F1

Remember that each class has its own precision, recall, and F1.

But Scikit-Learn requires that the `scoring=` parameter be a single number.

For this, we can use

- "precision\_macro"
- "recall\_macro"
- "f1\_macro"

which average the score over the classes.

*Example:*

```
cross_val_score(  
    pipeline,  
    X=X_train, y=y_train,  
    scoring="f1_macro",  
    cv=10  
) .mean()
```

0.6475574801118277

Kısaca:

Tek bir skor lazım → tüm sınıfların skorlarını ortalıyoruz (macro). Böylece GridSearchCV veya cross-validation gibi yöntemlerde kullanabiliriz.

## Precision-Recall Curve

Another way to illustrate the tradeoff between precision and recall is to graph the **precision-recall curve**.

## Precision-Recall Curve

Another way to illustrate the tradeoff between precision and recall is to graph the **precision-recall curve**.

First, we need the predicted probabilities.

## Precision-Recall Curve

Another way to illustrate the tradeoff between precision and recall is to graph the **precision-recall curve**.

First, we need the predicted probabilities.

```
y_train_probs_ = pipeline.predict_proba(X_train)  
y_train_probs_
```

```
array([[1. , 0. ],  
       [1. , 0. ],  
       [0.6, 0.4],  
       ...,  
       [1. , 0. ],  
       [0.8, 0.2],  
       [1. , 0. ]])
```

## Precision-Recall Curve

Another way to illustrate the tradeoff between precision and recall is to graph the **precision-recall curve**.

First, we need the predicted probabilities.

```
y_train_probs_ = pipeline.predict_proba(X_train)  
y_train_probs_
```

```
array([[1. , 0. ],  
       [1. , 0. ],  
       [0.6, 0.4],  
       ...,  
       [1. , 0. ],  
       [0.8, 0.2],  
       [1. , 0. ]])
```

By default, Scikit-Learn classifies a transaction as fraud if this probability is > 0.5.



## Precision-Recall Curve

Another way to illustrate the tradeoff between precision and recall is to graph the **precision-recall curve**.

First, we need the predicted probabilities.

```
y_train_probs_ = pipeline.predict_proba(X_train)  
y_train_probs_
```

```
array([[1. , 0. ],  
       [1. , 0. ],  
       [0.6, 0.4],  
       ...,  
       [1. , 0. ],  
       [0.8, 0.2],  
       [1. , 0. ]])
```

By default, Scikit-Learn classifies a transaction as fraud if this probability is  $> 0.5$ .



What if we instead used a threshold  $t$  other than 0.5?

## Precision-Recall Curve

Another way to illustrate the tradeoff between precision and recall is to graph the **precision-recall curve**.

First, we need the predicted probabilities.

```
y_train_probs_ = pipeline.predict_proba(X_train)  
y_train_probs_
```

```
array([[1. , 0. ],  
       [1. , 0. ],  
       [0.6, 0.4],  
       ...,  
       [1. , 0. ],  
       [0.8, 0.2],  
       [1. , 0. ]])
```

By default, Scikit-Learn classifies a transaction as fraud if this probability is  $> 0.5$ .



What if we instead used a threshold  $t$  other than 0.5?

Depending on which  $t$  we pick, we'll get a different precision and recall. We can graph this tradeoff.

# Precision-Recall Curve

Let's graph the precision-recall curve in a Colab.

